


# SimpleScalar MP

multiprocess simulator



*developed by* **Pau Erola**

*under guidance of* **Carles Aliagas**

*June 2006*

Departament d'Enginyeria

**[DΣIM]**

**Informàtica i  
Matemàtiques**



UNIVERSITAT  
ROVIRA I VIRGILI


# What's in this document

- Previous concepts and goals
- Specifications
- Design and development
- User's Guide and Simulation Examples
- Conclusions
- References

# Computer Evolution and Performance

The chart displays the theoretical maximum performance of various microprocessors over time. The y-axis represents the performance in million operations per second (logarithmic scale), and the x-axis represents the year.

Year	Processor	Performance (million operations per second)
1974	Intel 8080	~0.6
1974	Rockwell 6801	~0.6
1984	TI 9900	~10,000
1984	Motorola 68000	~10,000
1994	Intel 486	~10,000
1994	Zilog Z8000	~10,000
1994	Digital VAX-4000	~10,000



# Computer Architecture Simulators

- Highly competitive market

- ☐ increasing demand of performance and reliability
- ☐ need to have suitable tools to be the first

- An architectural simulator is a tool that reproduces the behavior of a computing device

- ☐ faster and more flexible development cycle
- ☐ permits more design space exploration
- ☐ facilitates validation before H/W becomes available

# SimpleScalar Tool Set

- Computer architecture research test bed
- Freely available with source and docs from UW-Madison
- First public release in July '96
- Primary Advantages
  - ☐ extensible
  - ☐ portable
  - ☐ detailed
  - ☐ performance

# Scheduling Processes

- Modern general-purpose computers have an operating system to run other programs
  - Multitasking operating systems share common processing resources for multiple tasks
    - cooperative multitasking
    - preemptive multitasking
  - The scheduler is the component of the kernel that selects which process to run next
    - multiprogramming systems
    - time-sharing systems
    - real-time systems





# Project Aims

- Multitasking operating systems produce lost of performance
  - Need to store and restore the CPU context to memory on context switch
  - Increment of cache accesses and misses
  - Decrement of branch predictor hits
- Our primary aim is to hack Sim-OutOrder to simulate multiple processes running concurrently using a configurable scheduling strategy



# Sim-OutOrder


- Sim-OutOrder is a very detailed performance simulator
  - generates timing statistics for a detailed out-of-order issue processor core with two-level cache memory hierarchy and main memory
- Source code design philosophy
  - infrastructure facilitates “rolling your own”
  - performance and flexibility before clarity



# Main code

```
ruu_init();  
for(;;)  
{  
    ruu_commit();  
    ruu_writeback();  
    lsq_refresh();  
    ruu_issue();  
    ruu_dispatch();  
    ruu_fetch();  
}
```

- in-order retirement
- wait for outputs
- execution
- wait for inputs and functional unit
- fetch and decode instr.

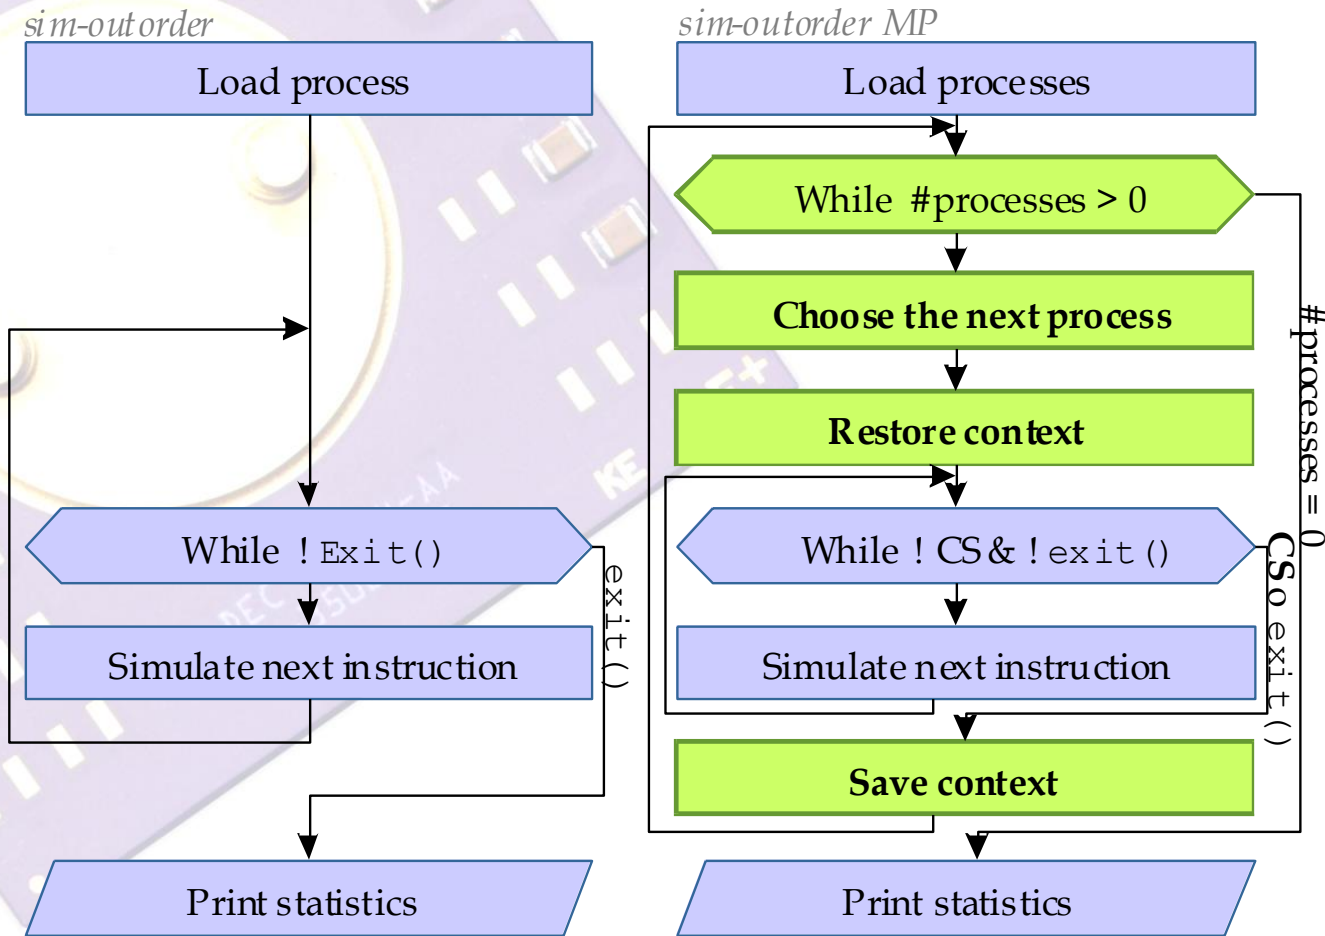




# Hacking Strategy

- Load all processes
- Add a configurable scheduler
  - ☐ Quantum
  - ☐ Flush of structures
  - ☐ Penalizations
  - ☐ Blocking
- Add new individual statistics for each process
- Do simulations to validate our modifications

# Execution Scheme

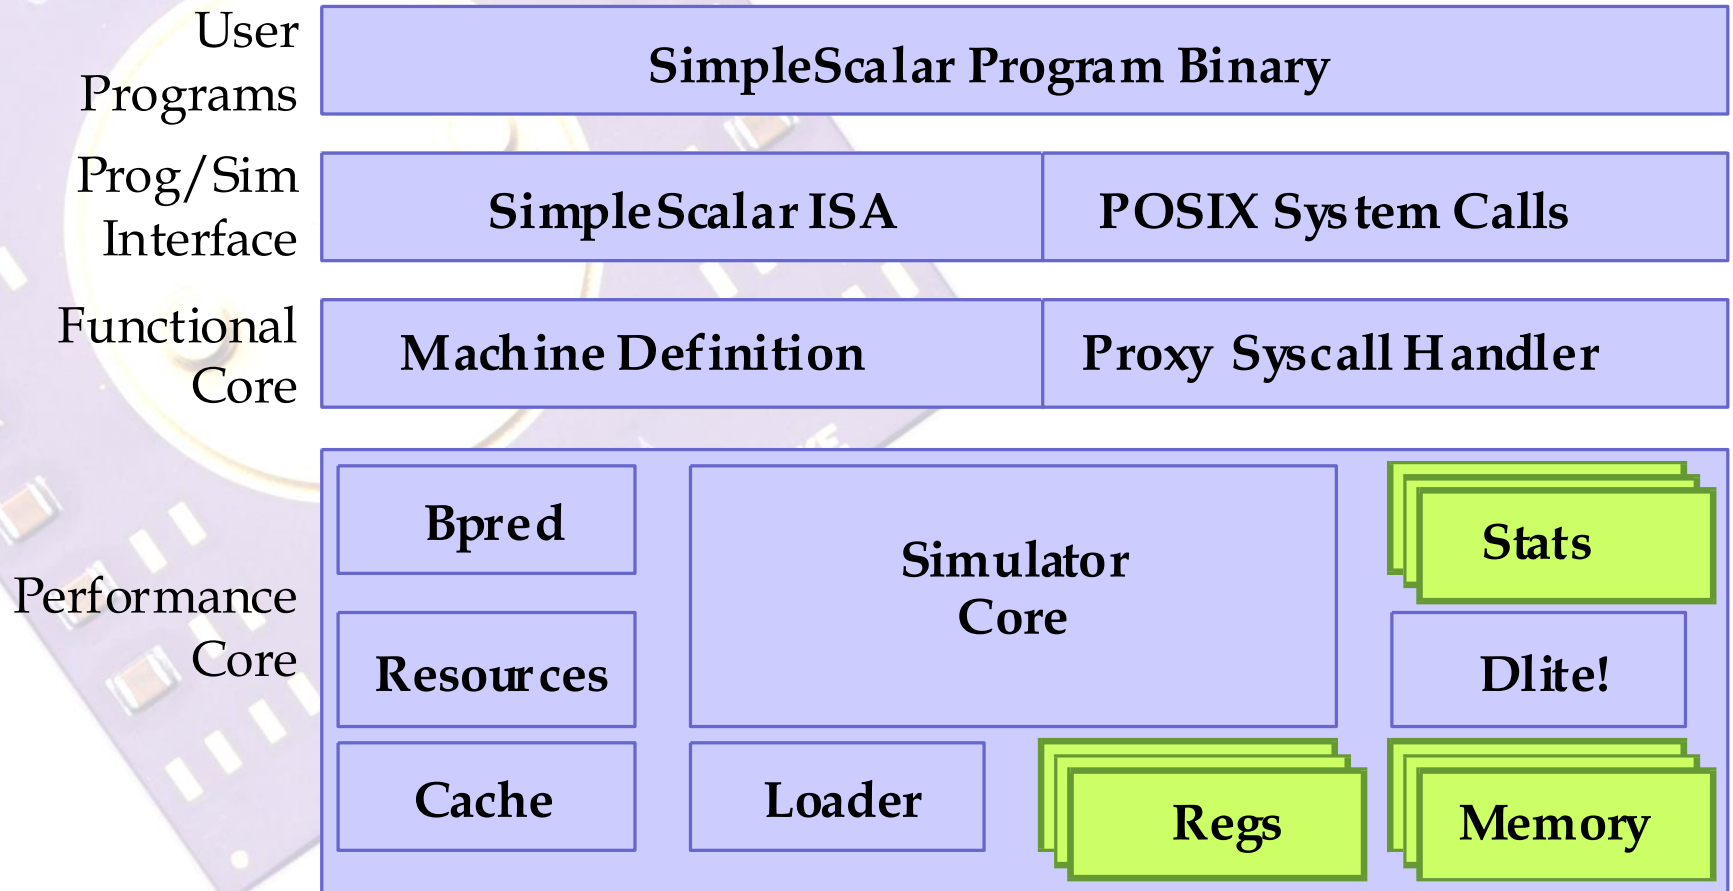




# Process Concept

- A process is a sequential program in execution
  - unique ID
  - memory space
    - object program to be executed
    - data on which the program will execute
  - virtual CPU
    - registers & PC
- Implemented in `schd.h`

# Simulator Structure



# Structure Implementation

## ■ Structure duplications

```
static struct regs_t mp_regs[LD_MAX_PROCS];  
static struct mem_t *mp_mem[LD_MAX_PROCS];  
mp_{nom_estadística} [LD_MAX_PROCS]
```

## ■ Memory identification

- SimpleScalar uses  $2^{33}$  bits of address
- With higher bits can identify the process' memory space

```
addr = addr | (ld_proc_id<<34);
```



# Memory Structure

Old Memory Space

1 0 0 0 0 0 0 0 0

33 bits address

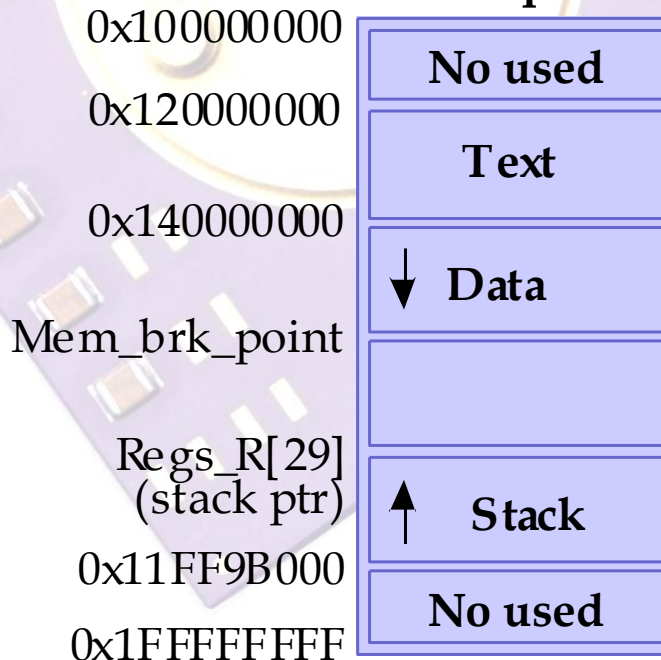
New Memory Space

OR PID

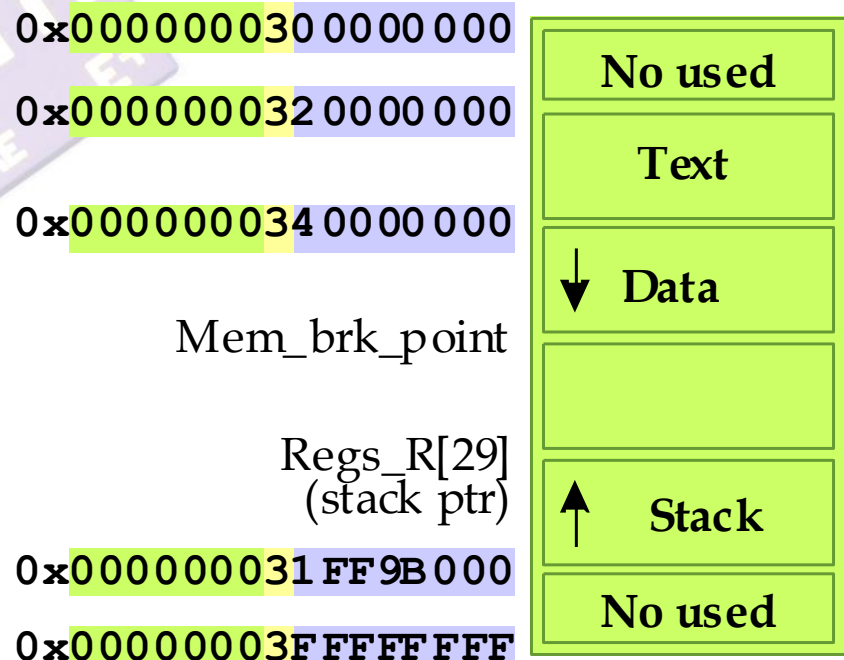
0 0 0 0 0 0 0 3 0 0 0 0 0 0 0 0

64 bits address

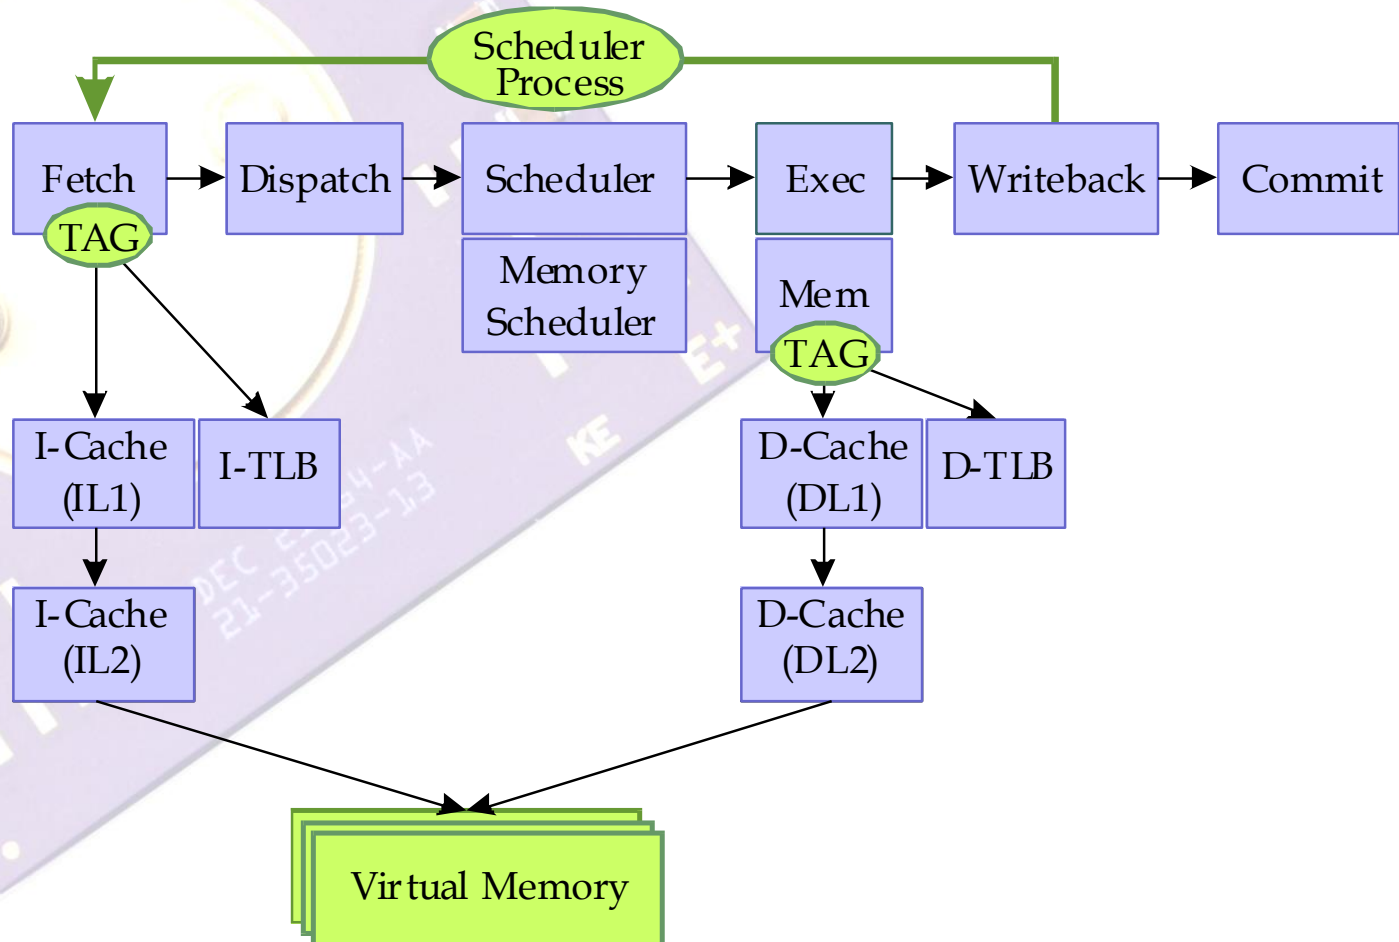
Unique



Process 1



# Multiprocess Out-of-Order Issue Simulator

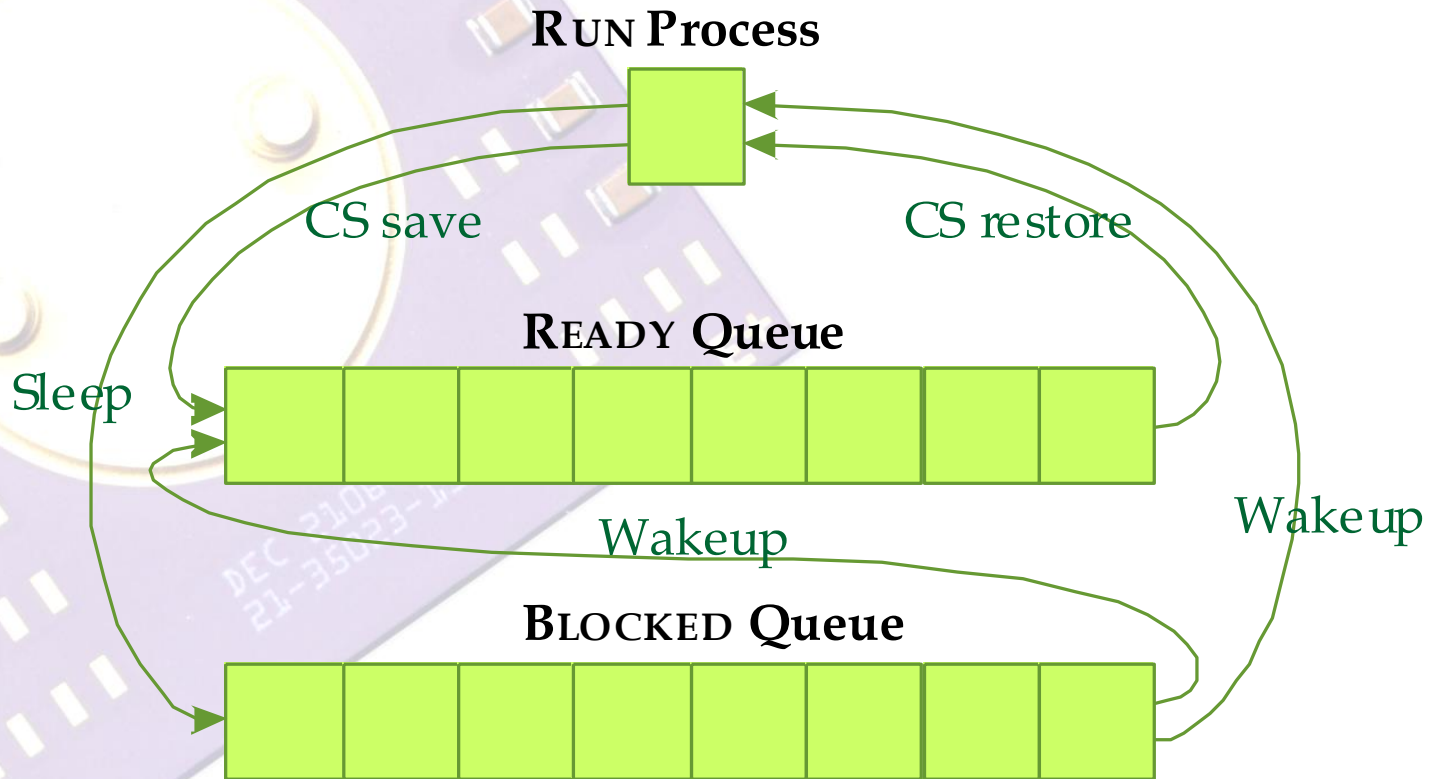


# Scheduler & Context Switch

- Implemented in sim-oomp.c

```
sim_cycle++  
wakeup_blocked_processes  
if ( context_switch_required ) then  
    context_save  
    select_next_process  
    context_restore  
endif
```

# Scheduler Queues



# Scheduler Queues Implementation

- RUN points to process that is currently executing  
`unsigned int ld_proc_id;`
- READY queue contains the processes that are waiting to run
  - queue ordered by priority (lower priority first)
- BLOCKED queue contains the processes that are waiting for a system call
  - queue ordered by wakeup cycle
- Macros for queues manipulation are implemented in `schd.h`

# Statistics Implementation

- Preserve the global simulation statistics
- Add new individual statistics for each process

`mp_{stat_name} [process_id]`



# User's Guide

## ■ Installation

```
$ gunzip $HOME/SimpleScalar-3.0-MP.tar.gz  
$ tar -vxf $HOME/SimpleScalar-3.0-MP.tar  
$ cd $HOME/SimpleScalar-3.0-MP  
$ make config-alpha  
$ make sim-oomp
```

## ■ Execution

```
$ sim-oomp <configs> list_of_executables
```

# List of Executables

## ■ Highly configurable simulator

**<quan>** quantum in cycles

**<penal>** cycles of penalization at CS

**<blk%>** probability of blocking during system calls

**<appr>** preemptive scheduler

**<pri>** priority of the process

**<wakeup>** start cycle

**<stdin>** **<exec>** **[arg]** stdin file, executable and arguments

**<dtlb>**, **<itlb>**, **<d11>**, **<d12>**, **<i11>**, **<i12>**  
flush structure at context switch

# List of Executables Example

```
# -s <quan> <penal> <blk%> <bskip> <appr> <inter>  
    <squan> <stdin> <sexec> [sargs]
```

```
-s 15000 150 60.0 false true false 200  
NULL $TESTS/test-printf
```

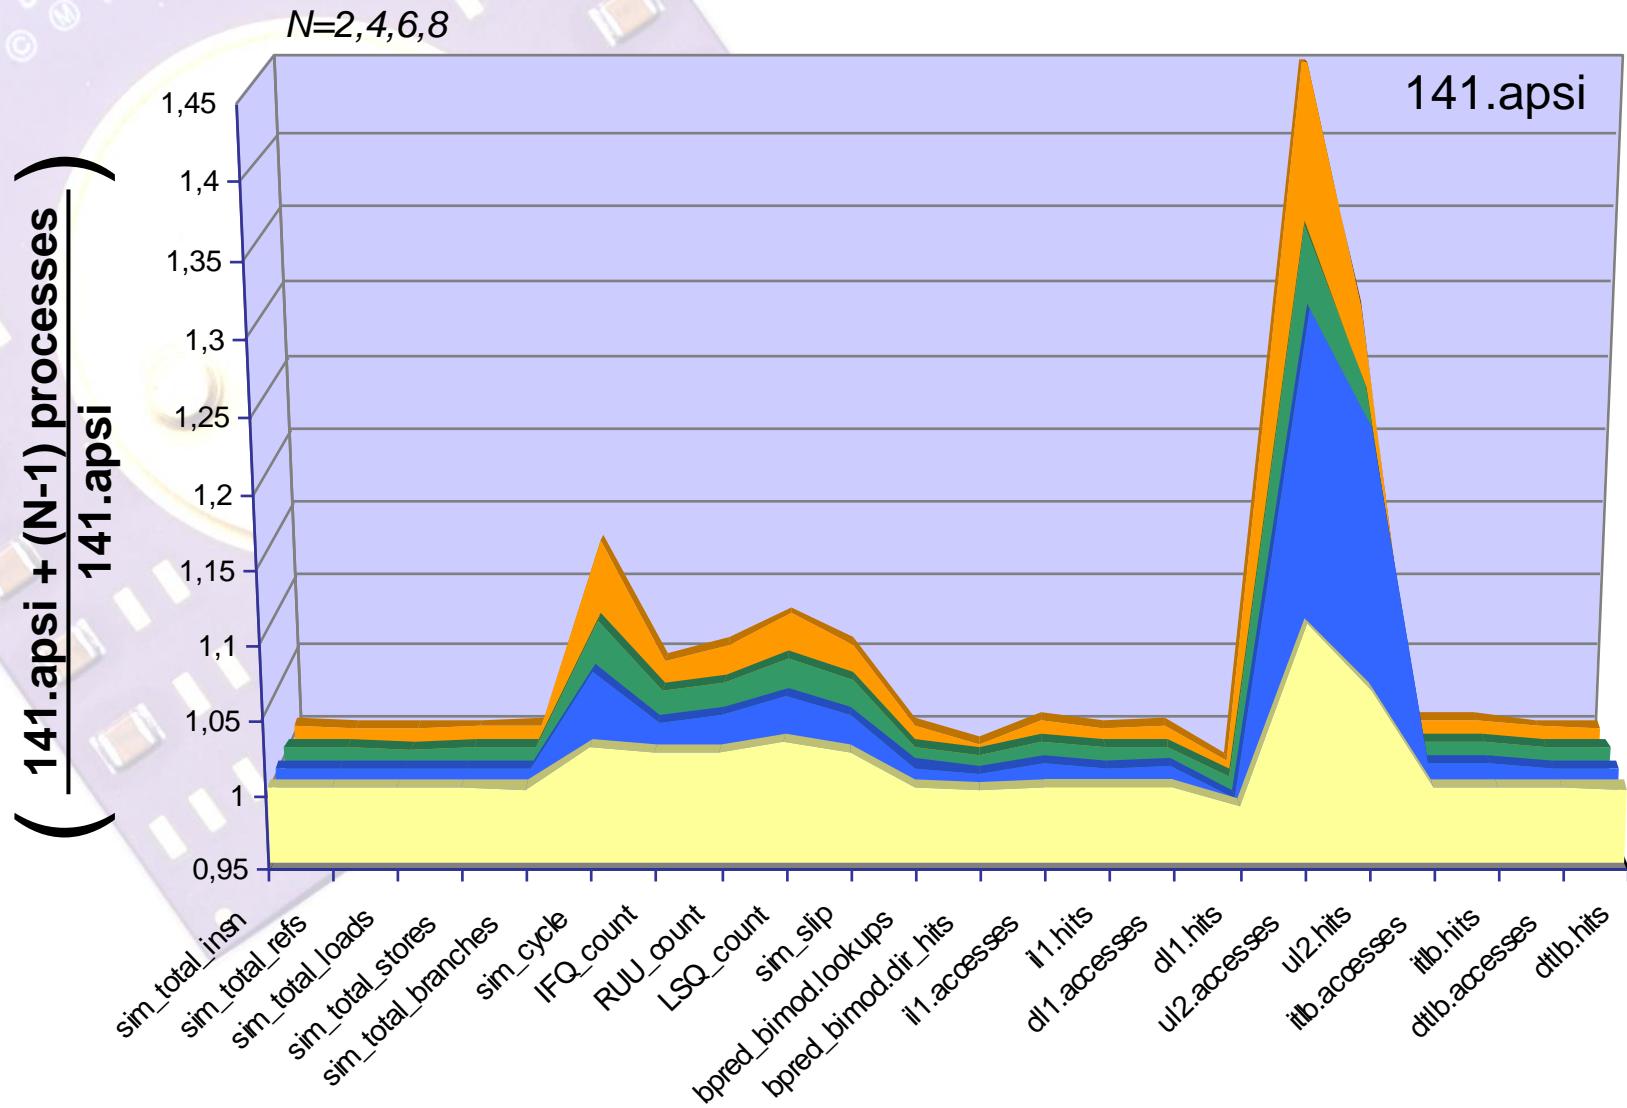
```
# -f <dtlb> <itlb> <d11> <d12> <i11> <i12>  
-f false false true false true false
```

```
# -e <pri> <exit> <wakeup> <fastfwd> <blk%> <maxins>  
    <stdin> <exec> [arg]
```

```
-e 2 s 6000 0 100.0 0 NULL $TESTS/test-  
args ARG1 ARG2
```

```
-e 5 t 0 500 20.0 0 NULL $TESTS/test-math
```

# Simulation of $N$ different processes




# Simulation of 8 processes with penalizations

( 8 proc. RR & preemption + penalization )  
8 proc. RR & preemption

■ + CS penalization

■ + CS penalization and blocking



# Interpretation of simulations

- The lost of performance in multitasking (more cycles) is reflected in
  - Higher L1 cache miss rate & more L2 cache accesses (higher latency)
  - Higher architecture queues (RUU, LSQ, IFQ) occupancy
  - More branch predictor accesses & higher TLB miss rate
- More concurrent processes produce less efficiency
- System overheads (scheduler and blocking) produce minimum lost of performance (previous simulation)





# Conclusions

- Multitasking operating systems produce lost of performance
  - Context switch overhead
  - Cache and speculative misses
- Research Topics
  - hardware CS approach
  - cache replacement algorithms
  - branch prediction algorithms
  - cache memory structure/configuration

# References

- T.M. Austin, "*A User's and Hacker's Guide to the SimpleScalar Architectural Research Tool Set (for tool set release 2.0)*", Intel MicroComputer Research Labs, gener 1997
- A. Silberschatz, P. Galvin, G. Gagne, *Sistemas Operativos (6th edition)*, Limusa Wiley, 2002
- J.S. Evans, R.H. Eckhouse, *Alpha<sup>TM</sup> RISC Architecture for programmers*, Prentice Hall, 1999
- [www.simplescalar.com](http://www.simplescalar.com)